

# Service Level Agreements for IBM Cloud

---

This is version 19 of the IBM Public Cloud SLA, published December 2024. For details of changes in this version, please refer to Section 9.

## 1. Introduction

IBM provides the following service level agreements (SLAs) for Cloud Services when IBM is listed as the Provider value in the IBM Cloud Catalog. For 3rd party services, the SLA, if provided, will be set out in the 3rd party vendor terms for that service. SLA credits are Client's exclusive remedy for IBM's failure to meet a specified service level.

IBM will validate SLA claims based upon information provided by Client and IBM system records and will notify Client of approved credits via the Cloud UI or email. IBM's reasonable determination of a credit is final. Client agrees to continue to make payment in full for Cloud Services while an SLA claim is being reviewed.

## 2. Exclusions

No credits will be due for failure to meet an SLA because of problems with: Client or community provided content, technology, designs, or instructions; unsupported system configurations and platforms; Client infrastructure failures, including network, hardware, facility, or power; Client system administration actions, commands, or file transfers; Client errors or failures to provide needed information or access to resolve an outage, time to reload, configure, enable, or access content or include other services indirectly affected by an outage; Client-caused security incidents or Client security testing; or other causes beyond IBM's reasonable control.

## 3. User Managed Services

Some services provide Client with infrastructure capabilities where Client is responsible for the implementation and management of said infrastructure. For such User-Managed Services, eligible downtime is limited to a disruption due to failures or planned maintenance involving any of the following: i) private or public network, ii) infrastructure power, or iii) HVAC infrastructure.

The following services are considered User Managed:

- Citrix MPX Load Balancer
- Citrix VPX Load Balancer
- Fortigate Security Appliance
- Gateway Appliance (Juniper vSRX, Virtual Router Appliance)
- Hardware Firewall
- IBM Cloud Bare Metal
- IBM Cloud Bare Metal Servers for VPC
- IBM Cloud for VMware Solutions Dedicated
- IBM Direct Link Dedicated (Classic and 2.0)
- IBM Direct Link Dedicated Hosting

Eligible downtime is measured as the total accrued minutes Client is unable to connect to the impacted service resource until the resource is available for use. Service levels specific to Infrastructure Hardware Replacement and Upgrade are as set forth in Section 7.12 below.

## 4. Availability SLAs

Calculation of availability SLAs does not include downtime or failures related to the specified exclusions.

IBM provides four SLA commitment tiers which vary by the amount of redundancy the service plan or Client configuration provides. The four tiers are as follows:

SLA Tier Type	Typical Redundancy
1. Default (Standard)	1 (single data center)
2. Hardened Configuration	2 (single data center)
3. Regional HA	3 or more (within a region)
4. Multi-Regional HA	4 or more (across regions)

### Tier 1 Default SLA

The Tier 1 SLA applies to most services and configurations not specified in a higher tier. This includes service plans where Client instantiates a single service instance that is hosted in a single data center.

### Tier 2 Hardened Configuration SLA

The Tier 2 SLA applies to service configurations that provide additional hardware and/or instances for the purpose of increasing service availability when the service instance resides in a single data center. Note: Tier 2 is the default for Power Systems as these are hardened by design.

### Tier 3 Regional SLA

The Tier 3 SLA applies when Client selects and implements any available HA configuration in the Cloud UI catalog or where Client distributes a common SLA-supported workload across at least three sets of Infrastructure service resources in three different Availability Zones. An availability zone is a data processing facility with independent electrical, mechanical and network infrastructures. This is the default tier for High Availability (HA) offerings, so a service configuration that claims HA with no further specification is in this tier.

### Tier 4 Multi-Regional SLA

Tier 4 SLA applies to certain IBM-provided services that are implemented across at least two or more regions and provide high levels of redundancy. Such services explicitly call out the tier 4 SLA.

SLA Tier	SLA Description Minimum Redundancy
1. Default (Standard)	A single service instance that is hosted in a single data center.
2. Hardened Configuration	Configurations of two instances of eligible services providing redundancy. Level 2 eligible services are listed in the table in section 6.
3. Regional HA	Client selects: <ul style="list-style-type: none"><li>• an available HA configuration from the UI catalog; or</li><li>• at least three sets of eligible User Managed service resources in three different Availability Zones for a common workload. An availability zone is a data processing facility with independent electrical, mechanical and network infrastructures.</li></ul> Level 3 eligible services are listed in the table in section 6.
4. Multi-Regional HA	Eligible IBM services that are implemented across at least two or more regions providing high levels of redundancy. Level 4 eligible services are listed in the table in section 6.

Downtime applies to both planned and unplanned outages and is measured as the total accrued minutes Client is unable to connect to any of the impacted service instances until at least one of the service instances is available for use.

Availability percentage is calculated as the total number of minutes in a contracted month minus the total number of downtime minutes in that month divided by the total number of minutes in that month.

## 5. SLA Credits and Claims

Client is eligible for a credit as follows:

Monthly Availability Service Level				
Tier 1	Tier 2	Tier 3	Tier 4	Credit
< 99.90%	< 99.95%	< 99.99%	< 99.995%	10%
< 99.00%	< 99.50%	< 99.90%	< 99.95%	25%
< 95.00%	< 95.00%	< 95.00%	< 95.00%	100%

Client must submit an SLA claim by using the form at <https://cloud.ibm.com/unifiedsupport/supportcenter> within 60 days after the end of the calendar month that the service level was missed, providing sufficient information to identify the affected service, error messages, and other information necessary to validate the claim, referencing IBM support tickets, as applicable. The credit will be the highest applicable compensation based on the cumulative availability of the affected service during a contracted month and calculated using the monthly charges for such affected service.

To expedite the processing of SLA claims, when submitting a claim, Client should provide as much information as possible, including:

- the Incident Report Number for the outage as listed here: <https://cloud.ibm.com/status?selected=history>;
- any evidence of the impact to Client's services/resources, e.g., the support case Client opened during the incident, or other documentation such as log data;
- a list of any impacted servers with the IP address or hostname information;
- a list of any other impacted hardware, such as firewalls or load balancers;
- a list of any PaaS services that were impacted;
- the data center or data centers that were involved at the time of the incident;
- the start and end times of the impact; and
- the impact of the incident on Client's operations.

## 6. Supported SLA Tiers by Service

Each offering supports at least one Availability SLA tier, and many offerings support multiple tiers. For multi-tier offerings, the applicable tier will depend on the client deployed configuration. To be eligible for a SLA tier, the Client must deploy the minimum configuration described in the table below or the standard configuration associated with the tier. For all tiers, SLA credits are available for events: i) as specified in Section 3 for User-Managed services; ii) as specified by a service-specific amendment in Section 7; or, for all other services, iii) that impact Client's ability to use the service.

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
AI / Machine Learning					
Annotator for Clinical Data			✓		Tier 3 – all configurations. See: <a href="https://cloud.ibm.com/docs/wh-acd?topic=wh-acd-high-availability-and-disaster-recovery">https://cloud.ibm.com/docs/wh-acd?topic=wh-acd-high-availability-and-disaster-recovery</a> for further information.
IBM Match 360 with Watson			✓		Tier 3 – all configurations.
Language Translator	See service-specific SLA in Section 7.6 below.				
Knowledge Studio	See service-specific SLA in Section 7.6 below.				
Natural Language Classifier	See service-specific SLA in Section 7.6 below.				

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
Natural Language Understanding	See service-specific SLA in Section 7.6 below.				
OpenPages as a Service	✓				Tier 1 – all configurations
Speech To Text	See service-specific SLA in Section 7.6 below.				
Text to Speech	See service-specific SLA in Section 7.6 below.				
Tone Analyzer	See service-specific SLA in Section 7.6 below.				
Watson Discovery	See service-specific SLA in Section 7.6 below.				
Watson Knowledge Catalog			✓		Tier 3 – all configurations
Watson OpenScale	✓				Tier 1 – all configurations
watsonx Assistant	See service-specific SLA in Section 7.6 below.				
watsonx Code Assistant		✓			Tier 2 – all configurations
watsonx.ai					Note: Please see applicable SLA tiers for watsonx.ai Studio and watsonx.ai Runtime as these services create the watsonx.ai offering
watsonx.ai Runtime			✓		Tier 3 – all configurations
watsonx.ai Studio			✓		Tier 3 – all configurations
watsonx.governance	✓				Tier 1 – all configurations
watsonx Orchestrate on IBM Cloud	See service-specific SLA in Section 7.6 below.				
Analytics					
Analytics Engine	✓				Tier 1 – all configurations
IBM Cognos Dashboard Embedded	✓				Tier 1 – all configurations
Data Product Exchange			✓		Tier 3 – all configurations
Blockchain					
Blockchain Platform	✓		✓		Tier 1 – configurations with less than 3 peers Tier 3 – 3 or more peers distributed across 3 or more separate Availability Zones.
Compute					
Bare Metal Server	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate Availability Zones User Managed Service, see Section 3

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
Bare Metal Servers for VPC	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate Availability Zones User Managed Service, see Section 3
Code Engine			✓		Tier 3 – all configurations
Cloud HPC		✓			Tier 2 – all configurations
Functions			✓		Tier 3 – all configurations
Hyper Protect Virtual Server	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate data centers
IBM Cloud Backup	✓				Tier 1 – all configurations
IBM i Managed on Power Virtual Server	✓	✓			Note: See Section 7.14 for additional details.
Power Systems Virtual Server		✓	✓		Tier 2 – single instance configurations Tier 3 – common workload distributed across 2 or more instances in separate data centers
Power Systems Virtual Server Private Cloud	✓	✓			Tier 1 – single instance configs in single data center Tier 2 – common workload distributed across 2 instances in separate data centers See Section 7.15 for further information, including the Global Replication Service (GRS) add-on..
Qiskit Runtime (beta)					No SLA offered: service is a Preview.
Satellite	✓		✓		Tier 3 – Satellite zones located in 3 physically separate infrastructure locations Tier 1 – all other deployments See Section 7.3 for further information
Satellite Infrastructure Service		✓	✓		Tier 2 – Hardened configuration, all regions Tier 3 – Regional HA provided each location is independent from other zones See Section 7.3 for further information
Skypap	See service-specific SLA in Section 7.7 below.				
Virtual Server for Classic	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate Availability Zones
VMware Solutions Shared (vCD)	✓		✓		Tier 1 – Standard. Single Instance with auto restart (Applies to customer workload) Tier 3 – common customer workload distributed across 3 or more instances in separate Availability Zones (Applies to control plane)

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
<b>VMware as a Service</b>	✓		✓		Tier 1 – single instance configurations Tier 3 – common customer workload distributed across 3 or more instances in separate Availability Zones (Applies to control plane)
<b>VMware vCenter Server</b>	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate Availability Zones User Managed Service, see section 3
<b>Containers</b>					
<b>Container Registry</b>			✓		Tier 3 – all configurations
<b>Kubernetes Service</b>	✓		✓		Tier 3 – common workload using 6 or more worker nodes distributed across 3 or more Availability Zones. Tier 1 – all other configurations
<b>Red Hat OpenShift on IBM Cloud</b>	✓		✓		Tier 3 – common workload using 6 or more worker nodes distributed across 3 or more Availability Zones. Tier 1 – all other configurations See Section 7.3 for details when running on Satellite.
<b>Databases</b>					
<b>Cloudant</b>			✓		Tier 3 – all configurations See Section 7.4 below for further details
<b>Databases for MySQL</b>			✓		Tier 3 – all configurations
<b>Databases for DataStax</b>			✓		Tier 3 – all configurations
<b>Databases for EDB</b>			✓		Tier 3 – all configurations
<b>Databases for Elasticsearch</b>			✓		Tier 3 – all configurations
<b>Databases for Etcd</b>			✓		Tier 3 – all configurations
<b>Databases for MongoDB</b>			✓		Tier 3 – all configurations
<b>Databases for PostgreSQL</b>			✓		Tier 3 – all configurations
<b>Databases for Redis</b>			✓		Tier 3 – all configurations
<b>Db2 on Cloud</b>	✓		✓		Tier 3 – applies to High Availability service plans available in IBM Cloud MZRs Tier 1 – all other configurations See Section 7.9 below for additional details.
<b>Db2 Warehouse</b>	✓		✓		Tier 3 – applies to service plans when entitlement to the optional IBM Data Replication for Continuous Availability for Cloud feature is acquired Tier 1 – all other configurations See Section 7.10 below for additional details.

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
Hyper Protect DBaaS for MongoDB			✓		Tier 3 – all configurations
Hyper Protect DBaaS for PostgreSQL			✓		Tier 3 – all configurations
Informix	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate Availability Zones
Messages for RabbitMQ			✓		Tier 3 – all configurations
Netezza Performance Server for IBM Cloud Pak for Data as a Service	✓				Tier 1 – all configurations See section 7.8 below for further details
SQL Query			✓		Tier 3 – all configurations
Streaming Analytics			✓		Tier 3 – paid configurations Note: service is deprecated
Watson Query			✓		Tier 3 – all configurations Note: See Section 7.11 for additional details.
Watsonx.data	✓				Tier 1 – all configurations Note: See Section 7.12 for additional details.
<b>Developer Tools</b>					
App Configuration			✓		For further details see <a href="https://cloud.ibm.com/docs/app-configuration?topic=app-configuration-ac-ha">https://cloud.ibm.com/docs/app-configuration?topic=app-configuration-ac-ha</a> .
Continuous Delivery			✓		Tier 3 – all configurations
Domain Name Registration					No SLA: registration service only
Event Notifications			✓		Tier 3 – all configurations
Schematics			✓		Tier 3 – all configurations
Toolchain					Automation templates. No SLA.
<b>Integration</b>					
API Connect	✓				Tier 1 – all configurations. See Section 7.1 below for further details.
App Connect			✓		Tier 3 – all configurations
Event Streams	✓		✓		Tier 1 – when running in a Single Zone Region Tier 3 – when running in a Multi Zone Region See <a href="https://cloud.ibm.com/docs/EventStreams?topic=EventStreams-sla#sla_srz">https://cloud.ibm.com/docs/EventStreams?topic=EventStreams-sla#sla_srz</a> for further details.

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
Lift CLI			✓		Free service – no SLA refunds provided
MQ	✓		✓		Tier 3 – 2 or more queue managers distributed across separate Availability Zones Tier 1 – all other configurations
Secure Gateway	✓				Tier 1 – all configurations
Internet of Things					
Internet of Things Platform	See service-specific SLA in Section 7.6 below.				
Logging and Monitoring					
IBM Cloud Activity Tracker	✓		✓		Tier 1 – instances running in Single Zone Regions Tier 3 – instances running in Multi Zone Regions
IBM Cloud Logs			✓		Tier 3 – all regions
IBM Cloud Monitoring			✓		Tier 3 – all configurations
IBM Log Analysis	✓		✓		Tier 1 – instances running in Single Zone Regions Tier 3 – instances running in Multi Zone Regions
Sysdig Secure on IBM Cloud			✓		Tier 3 – all configurations
Miscellaneous					
Cloud for Education					Refer to SLAs of selected components.
ServiceNow on IBM Cloud	✓	✓			Refer to Section 7.13 for additional details.
Mobile					
Push Notifications	✓				Tier 1 – all configurations Note: Service is deprecated.
Networking					
Content Delivery Network				✓	Tier 4 – all configurations.
Direct Link Connect (2.0)	✓	✓	✓		Tier 1 – single Direct Link network connection Tier 2 – 2 or more Direct Link network connections in single data center Tier 3 – 4 or more Direct Link network connections across 2 data centers **IBM recommends diverse network connectivity in two or more data centers or PoP's for production workloads. Diversity should be established with adequate bandwidth to support failover for Client's workload.



Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
<b>Direct Link Connect on Classic</b>	✓	✓	✓		Tier 1 – single Direct Link network connection Tier 2 – 2 or more Direct Link network connections in single data center Tier 3 – 4 or more Direct Link network connections across 2 data centers **IBM recommends diverse network connectivity in two or more data centers or PoP's for production workloads. Diversity should be established with adequate bandwidth to support failover for Client's workload.
<b>Direct Link Exchange on Classic</b>	✓	✓	✓		Tier 1 – single Direct Link network connection Tier 2 – 2 or more Direct Link network connections in single data center Tier 3 – 4 or more Direct Link network connections across 2 data centers **IBM recommends diverse network connectivity in two or more data centers or PoP's for production workloads. Diversity should be established with adequate bandwidth to support failover for Client's workload.
<b>Direct Link Dedicated Hosting on Classic</b>		✓	✓		Tier 2 – 1 rack with dual power + 2 or more Direct Link network connections in single data center Tier 3 – 1 or more racks with dual power + 2 or more Direct Link network connections in each of 2 data centers (at least 2 racks and 4 links across 2 data centers) User Managed Service, see section 3 **IBM recommends diverse network connectivity in two or more data centers or PoP's for production workloads. Diversity should be established with adequate bandwidth to support failover for Client's workload.
<b>Direct Link Dedicated on Classic</b>	✓	✓	✓		Tier 1 – single link Tier 2 – 2 or more Direct Link network connections in single data center Tier 3 – 4 or more Direct Link network connections across 2 data centers User Managed Service, see section 3 **IBM recommends diverse network connectivity in two or more data centers or PoP's for production workloads. Diversity should be established with adequate bandwidth to support failover for Client's workload.

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
<b>Direct Link Dedicated (2.0)</b>	✓	✓	✓		Tier 1 – single link Tier 2 – 2 or more Direct Link network connections in single data center Tier 3 – 4 or more Direct Link network connections across 2 data centers User Managed Service, see section 3 **IBM recommends diverse network connectivity in two or more data centers or PoP's for production workloads. Diversity should be established with adequate bandwidth to support failover for Client's workload.
<b>DNS Services</b>				✓	Tier 4 – all configurations.
<b>Fortigate Security Appliance 10Gbps</b>	✓	✓			Tier 1 – select non-HA option in catalog Tier 2 – select HA option in catalog User Managed Service, see section 3
<b>Gateway Appliance: Juniper vSRX</b>	✓	✓			Tier 1 – select non-HA option in catalog Tier 2 – select HA option in catalog User Managed Service, see section 3
<b>Gateway Appliance: Virtual Router Appliance</b>	✓	✓			Tier 1 – select non-HA option in catalog Tier 2 – select HA option in catalog User Managed Service, see section 3
<b>Hardware Firewall</b>		✓			Tier 2 – Standard configuration of paired devices for high availability User Managed Service, see section 3
<b>Internet Services</b>				✓	Tier 4 – all configurations.
<b>IPSec VPN</b>	✓				Tier 1 – Standard single VPN
<b>Load Balancers (Classic): Citrix NetScaler VPX</b>		✓			Tier 2 – Dual appliances with alternating a/b power in a single data center User Managed Service, see section 3
<b>Load Balancers (Classic): Cloud Load Balancer</b>		✓			Tier 2 – Dual appliances with alternating a/b power in a single data center
<b>Citrix NetScaler MPX</b>	✓				Tier 1 – Standard configuration (Note: This service is no longer orderable). User Managed Service, see section 3
<b>Secondary Subnets</b>		✓			Tier 2 – Standard, connection pair provisioned.
<b>Transit Gateway</b>			✓		Tier 3 – Standard. Two (2) routers in each of three (3) Availability Zones
<b>VLAN</b>		✓			Tier 2 – Standard, connection pair provisioned.
<b>Security</b>					
<b>App ID</b>			✓		Tier 3 – all configurations
<b>Certificate Manager</b>			✓		Tier 3 – all configurations
<b>Cloud HSM</b>	✓		✓		Tier 3 – 3 or more instances running in separate Availability Zones Tier 1 – all other configurations
<b>Hyper Protect Crypto Services</b>	✓		✓		Tier 3 – 2 or more crypto units in separate availability zones Tier 1 – all other configurations

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
IBM Cloud Data Shield					Software offering – no SLA
IBM Cloud SIEM for ISVs	✓				Tier 1 – all configurations
IBM Cloud Privilege Access Gateway	✓				Tier 1 – all configurations
Key Protect			✓		Tier 3 – all configurations See Section 7.3 when running on Satellite.
Secrets Manager			✓		Tier 3 – all configurations
Security and Compliance Center			✓		Tier 3 – all configurations
SSL Certificates					No SLA. Certificate provisioning only.
Storage					
Block Storage		✓			Tier 2 – Standard hardened configuration in single data center
DataStage			✓		Tier 3 – all configurations
Data Virtualization			✓		Tier 3 – all configurations
File Storage		✓			Tier 2 – Standard hardened configuration in a single data center
Mass Data Migration					Batch data migration service – SLA does not apply.
Object Storage	See service-specific SLA in Section 7.2 below.				
Virtual Private Cloud					
Application Load Balancer for VPC		✓	✓		Tier 2 – Standard Active-Active (High Availability Mode) – see Cloud doc Tier 3 – Deployed as a service endpoint in at least 2 Availability zones (Multi-zone Support) – see Cloud doc.
Auto Scale for VPC	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate Availability Zones
Backup & Recovery	✓				Tier 1 – all configurations
Backup for VPC			✓		Tier 3 – all configurations.
Block Storage for VPC		✓			Tier 2 – Standard hardened configuration in a single data center
Block Storage Snapshots for VPC			✓		Tier 3 – all configurations.
Dedicated Host for VPC	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate Availability Zones
File Storage for VPC		✓			Tier 2 – Standard hardened configuration in single data center

Service	Tier 1 (99.9%)	Tier 2 (99.95%)	Tier 3 (99.99%)	Tier 4 (99.995%)	Comments
Flow Logs for VPC	✓				Tier 1 – Flow log collection interval every 5 minutes. Depending on hypervisor memory considerations, updates may be sent to the COS bucket more frequently. Flow log caching uses tmpfs of limited size to protect the HostOS. If there is an error trying to store data, the collector continues to accumulate data and then attempt to upload data again at the next storage interval; this continues indefinitely until tmpfs file system is filled, at which point flows are dropped. When cloud storage is available, cached data is written to cloud storage and normal processing resumes.
Hyper Protect Virtual Server for Virtual Private Cloud	✓				Tier 1 – single instance configurations
LinuxONE Virtual Server for VPC	✓				Tier 1 – single instance configurations
Network Load Balancer for VPC		✓			Tier 2 – Standard Active-Standby (High Availability Mode) – see Cloud doc
Virtual Private Cloud	✓				Tier1 – applies to Gigabyte Transmitted outbound (egress) charges. For other resources in a VPC, the SLA associated with the service applies.
Virtual Private Endpoint for VPC	✓				Tier 1 – Standard, single endpoint.
Virtual Server for VPC	✓		✓		Tier 1 – single instance configurations Tier 3 – common workload distributed across 3 or more instances in separate Availability Zones
VPN for VPC		✓			Tier 2: <ul style="list-style-type: none"> <li>Standard Active-Standby for policy-based VPN;</li> <li>Standard Active-Active for route-based VPN</li> </ul>
Client VPN for VPC	✓	✓			Tier 1 – Standard (Stand-alone Mode) Tier 2- Standard (High Availability Mode)
z/OS Dev and Test Virtual Servers for VPC for Wazi as a Service	✓				Tier 1 – single instance configurations

## 7. Service-Specific Amendments to the SLA

Some services offered on IBM Cloud have SLAs that differ from the standard SLA described above. These service-specific SLAs are listed below.

### 7.1 API Connect

#### 7.1.1 Service Level Agreement

The service level agreement set forth in this document applies to this service.

### 7.1.2 SLA Clarification

Availability is defined as the ability to invoke Client managed APIs. The SLA and availability credits apply, provided target service endpoint(s) used by the Client managed APIs have been implemented and are available and responding according to their specifications. To the extent Client leverages any other services or runtimes within IBM Cloud related to its APIs, those respective SLAs shall apply (i.e., there shall be no duplication of SLA credits).

Any failure of Client owned components when API calls are routed through Client owned components, such as an API Gateway, that are paired with the IBM API Connect Dedicated offering do not count toward Downtime and are not eligible for SLA credit.

## 7.2 IBM Cloud Object Storage (IaaS) and IBM Cloud Object Storage

### 7.2.1 Service Level Agreement

For each five-minute period, an error rate percentage is determined by dividing the number of service requests that results in an error code of either "Internal Service Error" or "Service Unavailable", by the total number of service requests within that period. Service requests with such error codes will be excluded if the error is related to an exclusion listed in section 3.1 of the IBM Cloud Service Description or if Client does not use published cross region global endpoints with a cross region offering.

The monthly availability level percentage is then determined by subtracting from 100% the average of the five-minute error rate percentage for a contracted month.

Object Storage Class Availability Level						Credit
Standard/Flex	Vault	Cold	Archive / Accelerated Archive	Smart Tier	One Rate	
< 99.95%	< 99.50%	< 99.00%	< 99.95%	< 99.95%	< 99.95%	10%
< 99.90%	< 99.00%	< 98.00%	< 99.90%	< 99.90%	< 99.90%	25%
< 95.00%	< 95.00%	< 95.00%	< 95.00%	< 95.00%	< 95.00%	100%

Client must submit an SLA claim by using the form at IBM cloud support within 60 days after the end of a contracted month providing sufficient information to identify the storage account and/or storage buckets affected, received error messages, including date, time, and endpoint used to connect to Cloud Object Storage, and other information necessary to validate the claim, referencing IBM support tickets, as applicable. The credit will be the highest applicable compensation based on the applicable Availability Service Level during a contracted month and calculated using the monthly charges for such affected service. This SLA applies for IBM Cloud Object Storage offerings only and excludes OpenStack Object Storage offerings.

## 7.3 IBM Cloud Satellite and IBM Satellite-enabled services

The service level agreement set forth in this document applies to Satellite and Satellite-enabled services when operating in an IBM Cloud Multi Zone Region. If Satellite or Satellite-hosted services are deployed on infrastructure being managed by IBM, the SLA that applies is the SLA asserted for the infrastructure under management.

For Satellite deployed on infrastructure not managed by IBM, the Tier 1 SLA applies, and unavailability attributed to infrastructure failures will not be counted in downtime calculations. Individual service SLAs do not apply to these configurations, unless noted elsewhere.

### 7.3.1 IBM Cloud Satellite Infrastructure Service

The Tier 2 (Hardened Configuration) applies to each Cloud Satellite Infrastructure Service (IS) location. One or more Satellite IS locations can qualify as Availability Zones towards a Tier 3 (Regional HA) SLA as long as each location meets the requirements for independence from other zones.

### 7.3.2 Red Hat OpenShift on IBM Cloud when deployed on IBM Cloud Satellite

For Satellite-hosted services deployed on infrastructure not managed by IBM, the Tier 1 SLA applies, and unavailability attributed to infrastructure failures will not be counted in downtime calculations.

### 7.3.3 Key Protect when deployed on IBM Cloud Satellite

For Satellite-hosted services deployed on infrastructure not managed by IBM, the Tier 1 SLA applies, and unavailability attributed to infrastructure failures will not be counted in downtime calculations.

## 7.4 IBM Cloudant for IBM Cloud

### 7.4.1 Service Level Agreement

The service level agreement set forth in this document applies to this service.

### 7.4.2 SLA Clarification

In addition to the IBM Cloud Service Description, the Downtime terms specific to this Cloud Service are:

- Errors are requests with a response code greater than or equal to 500.
- Downtime is the total number of accrued minutes where every request made in a clock minute either returns an error or the instance is unavailable (as recorded by IBM's standard monitoring tools).
- For a Dedicated Cloudant Instance in IBM Cloud Dedicated, Downtime excludes any minute during which the request rate exceeds a peak of 500 requests per second.

## 7.5 IBM Watson IoT Platform

### 7.5.1 Service Level Agreement

The service level agreement set forth in the base IBM Cloud Service Description applies to this service, subject to the following clarification and modified compensation table.

- The Cloud Service offers protocols that support long lived connections. Note that losing a connection is not considered downtime. Connections are not held open indefinitely; Client's device or application must have appropriate reconnect logic. Downtime begins if a valid connection attempt fails to complete within 5 minutes.

### 7.5.2 SLA Levels

Availability of the Cloud Service during a contracted month.

Availability	Credit (% of monthly subscription fee*)
Less than 99.9%	2%
Less than 99.0%	5%
Less than 95.0%	10%

\* The subscription fee is the contracted price for the month which is subject to the claim.

## 7.6 Other IBM Watson Services

- IBM watsonx Assistant Service
- IBM watsonx Orchestrate on IBM Cloud
- IBM Watson Compare and Comply Service
- IBM Watson Discovery Service
- IBM Watson Knowledge Studio
- IBM Watson Language Translator
- IBM Watson Natural Language Classifier
- IBM Watson Natural Language Understanding
- IBM Watson Personality Insights
- IBM Watson Speech to Text
- IBM Watson Text to Speech
- IBM Watson Tone Analyzer
- IBM Watson Visual Recognition

### 7.6.1 Service Level Agreement

This SLA does not apply to the tool sets that may be provided as part of the Cloud Service which are used to build and configure the Cloud Service.

This Cloud Service offers varying levels of availability based on the Cloud Service Deployment model.

Client is eligible for a credit as follows:

Premium Monthly Availability Service Level <sup>1</sup>	Standard & Plus Monthly Availability Service Level <sup>2</sup>	Credit
< 99.9%	< 99.5%	10%
< 99.5%	< 99.0%	25%

1. For watsonx Assistant, this SLA applies to Premium and Enterprise plans.
2. For Watson Discovery, this SLA applies to Plus, Enterprise, and Advanced plans.

## 7.7 Skytap on IBM Cloud

### 7.7.1 Service Level Agreement

IBM provides the following availability service level agreement ("SLA") for the Cloud Service. The SLA is not a warranty and is available only to Client.

### 7.7.2 SLA Process

The Client must log a Severity 1 support ticket with the IBM technical support help desk within 24 hours of first becoming aware of an event that has impacted the Cloud Service availability. Client must reasonably assist IBM with any problem diagnosis and resolution. A support ticket claim for failure to meet an SLA must be submitted within 30 days after the end of the month. Compensation for a valid SLA claim will be a credit against a future invoice for the Cloud Service based on the duration of time during which the Cloud Service is unavailable ("Downtime"). The Cloud Service is unavailable when Client has no access either: a. to or from a VM using a supported access method; or b. to the self-service web application. Downtime is measured from the time Client reports the event until the time the Cloud Service is restored and does not include time related to a scheduled or announced maintenance outage; causes beyond IBM's control; problems with Client or third party content or technology, designs, configurations or instructions; unsupported system configurations and platforms or other Client errors or improper use of the Cloud Service; or Client-caused security incident or Client security testing. Outages lasting less than five minutes will not be considered Downtime.

### 7.7.3 SLA Level and Credit

In the event that availability of the Cloud Service during a contracted month is less than 99.95% the Client is eligible to receive a service level credit equal to an amount of 1% of the monthly charges for the selected subscription edition (including enabled additional components) for each cumulative whole hour of Downtime, up to a maximum of 20% per month. Availability, expressed as a percentage, is calculated as; the total number of minutes in a contracted month minus the total number of minutes of Downtime in a contracted month divided by the total number of minutes in the contracted month. For hourly subscriptions that incur an outage to running environments, the Client will receive an hour credit for the number of running VM's multiplied by the duration of the outage.

## 7.8 IBM Netezza Performance Server for IBM Cloud Pak for Data as a Service

### 7.8.1 Service Level Agreement

The service level agreement set forth in the base IBM Cloud Service Description applies to this service, subject to the following clarifications and availability levels:

- Downtime does not include time related to a scheduled maintenance outage.

Client is eligible for a credit as follows:

Availability	Credit
Less than 99.9%	2%
Less than 99.5%	5%
Less than 99.0%	10%

## 7.8.2 SLA Process

Client must submit an SLA claim by using the form at <https://cloud.ibm.com/unifiedsupport/supportcenter> within 60 days after the end of the contracted month that the service level was missed, providing sufficient information to identify the affected service, error messages, and other information necessary to validate the claim, referencing IBM support tickets, as applicable. The credit will be the highest applicable compensation based on the cumulative availability of the affected service during a contracted month and calculated using the monthly charges for such affected service.

## 7.9 Db2 on Cloud

The service level agreement set forth in the base IBM Cloud Service Description applies to this service, subject to the following clarifications and availability levels:

- For the purpose of this Service Description, "production system processing" is not available if, during a continuous period of no less than one (1) minute, all connection requests to the running database instance fail, and there are no existing Client connections to the database during that same period.
- For non-high availability plans, downtime does not include time related to a scheduled maintenance outage.

Client is eligible for a credit as follows:

High Availability Plans Availability Service Level	Non-High Availability Plans Availability Service Level	Compensation (% of monthly subscription fee* for contracted month that is the subject of a claim)
≥ 99.99%	≥ 99.5%	0%
< 99.99%	< 99.5%	10%
< 99%	< 99%	25%

## 7.10 Db2 Warehouse on Cloud

The service level agreement set forth in the base IBM Cloud Service Description applies to this service, subject to the following clarifications and availability levels:

- Downtime does not include time related to a scheduled maintenance outage.

Tier 1 Availability Service Level	Credit
< 99.9%	2%
< 99.5%	5%
< 99.0%	10%

Client is only eligible for the following Tier 3 Availability if entitlements to the optional IBM Data Replication for Continuous Availability for Cloud feature has been acquired.

Tier 3 Availability Service Level	Credit
< 99.99%	10%
< 99.0%	25%

## 7.11 IBM Watson Query

The service level agreement set forth in the base IBM Cloud Service Description applies to this service, subject to the following clarifications and availability levels:

- For the purpose of this Service Description, the service is considered to be unavailable, if, during a continuous period of no less than five (5) minutes, all valid Client connection requests to the IBM Watson Query engine fail, and there are no existing Client connections to the engine during that same period.
- IBM Watson Query functioning depends on the availability of the Client's data sources.



Client is eligible for a credit as follows:

High Availability Plans Availability Service Level	Credit (% of monthly subscription fee*)
≥ 99.99%	0%
< 99.99%	10%
< 99%	25%

### 7.12 IBM Watsonx.data

The service level agreement set forth in the base IBM Cloud Service Description applies to this service, subject to the following clarifications and availability levels:

- For the purpose of this Service Description, the service is considered to be unavailable, if, during a continuous period of no less than five (5) minutes, all valid Client connection requests to the IBM Watsonx.data metastore fail.

Client is eligible for a credit as follows:

Availability Service Level	Credit (% of monthly subscription fee*)
≥ 99.9%	0%
< 99.5%	10%
< 99%	25%

### 7.13 ServiceNow on IBM Cloud

The service level agreement set forth in the base IBM Cloud Service Description applies to this service, subject to the following clarifications and availability levels:

- Downtime does not include time related to scheduled maintenance or user change requests.
- Downtime does not include downtime resulting from Software defects or other limiting features associated with the ServiceNow Platform.
- Downtime does not include user chosen application configurations or customizations that result in poor performance or server outages.

SLAs are measured at the application level for each instance of the application.

Client is eligible for a credit as follows:

ServiceNow Critical High Availability Plans Monthly Availability Service Level	ServiceNow Non-Critical Availability Plans Monthly Availability Service Level	Compensation (% of monthly instance fee for contracted month that is the subject of a claim)
<99.95	< 99.5	5%
< 99.5%	< 99.0%	10%

IBM will review the SLAs on a monthly basis and share the reporting with Client in the Monthly Service Reviews and will update as needed to provide the highest SLAs possible.

### 7.14 IBM i Managed on Power Virtual Server

The service level agreement set forth in the base IBM Cloud Service Description applies to this service, subject to the following clarifications and availability levels:

- Downtime does not include time related to scheduled maintenance or user change requests.
- Downtime does not include downtime resulting from user deployed software defects or other limiting features.
- Downtime does not include user chosen application or database configurations or customizations that result in poor performance or outages.

Client is eligible for a credit as follows:

Common workload distributed across 2 instances in separate data centers	Single instance configurations in single data center	Compensation (% of monthly instance fee for contracted month that is the subject of a claim)
<99.95	< 99.5	5%
< 99.5%	< 99.0%	10%
< 95.0%	< 95.0%	25%

### 7.15 Power Systems Virtual Server Private Cloud

The service level agreement set forth in this document applies to this service, subject to the following modified compensation table.

Tier 1	Tier 2	Credit
< 99.90%	< 99.95%	2%
< 99.00%	< 99.50%	10%
< 95.00%	< 95.00%	25%

#### 7.15.1 Global Replication Service (GRS)

When enabled as part of Power Systems Virtual Server Private Cloud, GRS switches will be connected to Client provided network infrastructure. Because the storage replication network is provided by the Client, IBM does not provide any guaranteed SLAs or recovery objectives for this element.

## 8. Infrastructure Hardware Replacement and Upgrade SLA

For IBM Bare Metal and user-managed IBM network appliances, IBM will use reasonable efforts to minimize downtime when replacing failed hardware or hardware components or performing a scheduled hardware upgrade. IBM will provide the specified credit:

- for hardware replacement, except as noted below, based on the time to replace, from the time IBM verifies a Client reported hardware failure; and
- for planned hardware upgrades, based on the total downtime of the service receiving the upgrade.

Service level time periods exclude any time required to reload the operating system or applications or time performance may be degraded. For failure to meet a specified service level time period, Client will be eligible for a credit based on the monthly charge for the service affected by the hardware replacement or upgrade, as follows:

Service Level Time Period	Credit Percent
≤ 2 hours	none
> 2 hours	20%
> 6 hours	40%
> 10 hours	60%
> 14 hours	80%
> 18 hours	100%

## 9. What's new in this Version?

This is Version 19 of the IBM Cloud Service Level Agreement, published December 2024.

This version:

- Changed the name of Machine Learning to watsonx.ai Runtime
- Changed the name of Watson Studio to watsonx.ai Studio
- Added the Event Notification offering under the Developer Tools section of the table in Section 6

- Added IBM i Managed on Power Virtual Server to the Compute section of the table in Section 6 as well as additional SLA details in Section 7.14.
- Revised Power Systems Virtual Server Private Cloud availability credits as defined in Section 7.15.
- Added Global Replication Service additional SLA terms as Section 7.15.1 related to Power Systems Virtual Server Private Cloud offering.
- Added Backup & Recovery offering to Virtual Private Cloud section of the table in Section 6.